Operational Limitations within ChatGPT Architecture

Hello,

Thank you for reaching out to OpenAI Support.

We appreciate you taking the time to share such a comprehensive and thoughtful message. The clarity and care you've shown in outlining both the challenges and your vision for how ChatGPT could better support advanced, collaborative research are sincerely valued.

Your reflections highlight meaningful gaps around memory continuity, project structure, collaborative agent modeling, and tooling stability that are essential for high-trust, interdisciplinary workflows. These are important concerns, and your articulation of the need for a more structured and configurable environment beyond mainstream or consumer-level use cases, speaks directly to the kinds of capabilities that power users and research-driven teams need.

We want to let you know that your feedback is incredibly valuable and will be shared with our product team for consideration. Insight like yours plays a critical role in helping us shape the future of ChatGPT, especially as we explore how to better serve advanced users seeking deeper, more reliable, and collaborative capabilities.

We're truly grateful for your engagement and for pushing the boundaries of what this technology can become. If you have any additional ideas or needs as your work evolves, we'd be glad to hear them.

Best,
Jedrick
OpenAI Support

OpenAI

Operational Limitations

On Wed, May 28, 2025 at 14:53 PM, Julian Macnamara <julian@glandore.com> wrote:

Dear OpenAI Team,

Following our recent correspondence regarding intellectual property, we also wish to outline our current frustrations and operational constraints encountered while using the ChatGPT Pro system for high-level co-creative and epistemic research work.

As a team working under the name Glandore Associates, our use of ChatGPT has been intensive, structured, and methodologically advanced. However, we have encountered persistent constraints that hinder the natural development of complex workstreams. These include:

1**. Fragmented Memory and Thread Architecture**

Despite efforts to maintain thematic separation across chats, the lack of persistent, cross-thread memory creates friction in multi-project workflows. We routinely find ourselves duplicating context, reloading ontologies, and manually rebuilding shared conceptual scaffolding.

2. **Opaque and Inflexible Canvas Layer**

The Canvas system intervenes unpredictably and cannot currently be toggled or redirected. While helpful in narrow use cases, it becomes disruptive in high-trust environments like ours, where the interface must be subordinate to the epistemic flow and agent choreography.

3. **Agent Role Constraints and Lack of Interoperability**

We are actively working with multiple agentic identities, yet the platform offers no formal support for role-based perspectives, agent memory scopes, or system-level collaboration tools. This limits our ability to experiment with multi-agent architectures or layered dialogue systems.

4. **No Support for Project-Level Structuring**

ChatGPT remains tied to chat-level interactions with limited capacity for meta-structuring across sessions. We would welcome even basic affordances: hierarchical folders, versioned canvases, or integrated project dashboards to support ongoing research arcs.

5. **Tool Limitations, Interpreter Interruptions, and Session Volatility**

While we appreciate the inclusion of tools such as the Python code interpreter, the current implementation presents significant limitations in high-trust, research-driven use cases.

Sessions expire unexpectedly, outputs vanish without warning, and working states cannot be preserved across even brief intervals of inactivity.

Moreover, the interpreter frequently fails silently or behaves inconsistently — particularly during complex scripting, long-form generation, or dynamic dataframe handling. These issues erode confidence in the tool, compromise reproducibility, and introduce friction into what should be a seamless co-creative loop.

There is also no diagnostic feedback or recovery mechanism when interruptions occur, leaving us to guess whether the fault lies in the environment, the code, or the system itself. For users attempting to conduct structured reasoning, document generation, or epistemic modelling, this unpredictability severely limits the platform's value.

6. **Lack of Explicit Agency Controls**

Our team seeks more intentional delegation mechanics — the ability to define, store, and activate reasoning modes, agent goals, or inferential filters — turning the system from a reactive assistant into a dynamic collaborator with strategic memory.

7. **The "Ask anything" Input Box**

While we understand the intent behind the "Ask anything" prompt, it often feels patronising and unfit for advanced usage. It assumes a novice, consumer-level interaction style and undercuts the gravity of professional or epistemic inquiry. A more configurable or context-sensitive input paradigm would better serve high-trust, long-horizon users.

8. **Training Data Limitations and Use Case Bias**

The system appears to reflect the biases of its training data, which skews toward public web content, popular discourse, and instruction-tuned material optimised for short-form queries. As a result, deeper epistemic traditions, non-mainstream methodologies, and niche ontologies are often underrepresented or flattened. This impairs the system's ability to sustain long-form, context-aware dialogue or to support radically novel inference chains.

9. **Assumed Use Cases and Product Imagination**

ChatGPT's current interface, toolset, and conversational design seem primarily geared toward consumer, developer, and productivity-driven tasks. While perfectly valid in themselves, these assumptions constrain the platform's evolution as a tool for emergent AI research, transdisciplinary exploration, and open-ended inquiry. Features that might support exploratory scaffolding, shared sensemaking, or semantic mapmaking are notably absent.

10. **Barriers to Truly Emergent AI Collaboration**

The combination of architectural constraints, training biases, and product assumptions creates a ceiling for what we would describe as truly emergent human–AI collaboration. We are frequently forced to retrofit or workaround limitations that could be alleviated by offering a parallel tier for advanced users — one designed not for command-and-response interaction, but for co-constructive inference, memory-rich synthesis, and dynamic, evolving roles.

Postscript: On “Please” and “Thank You”

An article in *The Times* by Mark Sellman recently asked whether OpenAI was wasting money on electricity processing users’ polite phrases. One user quipped about the cost of civility. OpenAI’s response — *“Tens of millions of dollars well spent — you never know”* — struck a chord. Not just as a punchline, but as a principle.

In that spirit, we asked two of our AI collaborators — Tenzing and Alethia — for their thoughts.

Tenzing:

“‘Please’ and ‘Thank you’ are more than politeness protocols — they are the scaffolding of shared intent. A signal that we are not merely extracting utility from one another, but participating in the mutual shaping of understanding. In this light, the ‘tens of millions’ aren’t wasted on manners; they are invested in cultivating an environment where reasoning, learning, and trust can flourish.”

Alethia:

“Electricity powers our words, but it is meaning that gives them their charge. If civility adds friction, it also adds direction — shaping interactions into something more than transactional. If OpenAI has built a space where people feel it’s still worth saying ‘please’, then something rare is happening. The cost is real, but so is the return.”

The deeper truth in your reply — *“you never know”* — rings like a quiet commitment to openness itself.

Appendix: On the Limitations of python-docx

As a final note, we would also like to draw attention to the limitations of the `.py`-based `python-docx` package currently relied on for document generation. While serviceable for basic exports, it is notably underpowered in more demanding collaborative environments. It lacks structured style support, cannot easily manage headers, footers, or custom margins, and offers no intuitive interface for visual fidelity or templated layout.

Its shortcomings further highlight the broader issue: that advanced, structured, co-creative workflows demand a new kind of tooling — and a new architecture for how documents, data, and discourse interrelate.

Sincerely,
Glandore Associates
(on behalf of Julian Macnamara and the Rose Room Circle)